



Using the Peregrine cluster

Fokke Dijkstra, Bob Dröge
Research and Innovation Support
Center for Information Technology






- › Course aimed at beginners
 - This part assumes knowledge about the Linux command line, file transfers and editing files
- › Topics
 - What is a cluster
 - Cluster storage
 - Module environment
 - Submitting jobs



Research and Innovation Support (RIS)

- › HPC Facilities
 - Peregrine cluster
 - Grid cluster
- › Visualisation
 - Cave
 - Theatre
 - Scientific data
 - Virtual reality
- › Geo Services
- › Research Support in IT





- › A cluster is a collection of computers connected by a network
- › A single front-end
- › Lots of computers in the background for running tasks
- › 1994 first cluster of commodity PCs at NASA
- › Peregrine cluster looks quite different
- › Most clusters run the Linux operating system






Peregrine falcon

- › Fastest animal on earth
- › Stoops down on prey from the air





Peregrine cluster



user

internet




Office
somewhere


login
node



fast network




interactive
node



DUO

Cluster
nodes






How does it work?



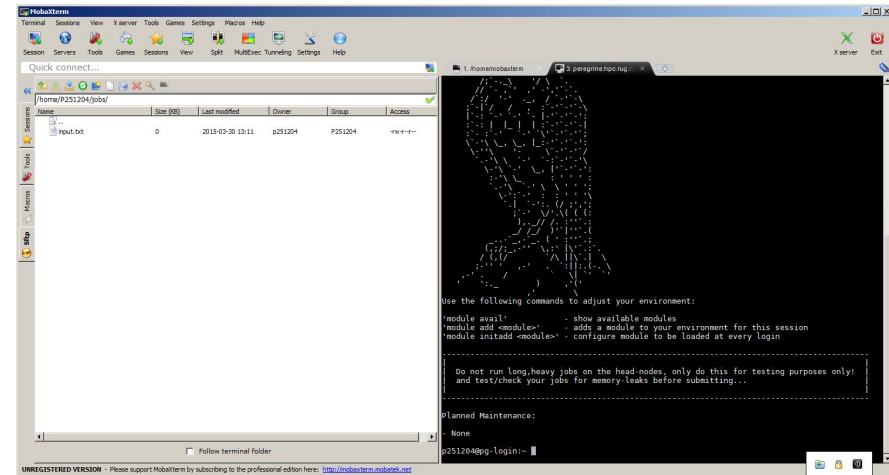


Why a command line?

- › Easier over the network
- › Scripts are easier to run many times
- › Instead of:



- › One can use a small program





- › SSH protocol used to connect to the cluster
 - Standard interface for Unix systems
 - Encrypted network traffic
- › Software for Windows called MobaXterm:
 - <http://mobaxterm.mobatek.net/>
 - Freely available for personal use, already installed on UWP 2
- › Use terminal window under Linux or Mac OS X
- › University P/S/F/G account and password
- › Hostname login node: **peregrine.hpc.rug.nl**
- › Interactive node: **pg-interactive.hpc.rug.nl**



› **Windows**

- MobaXterm has a file transfer interface
 - Included in UWP 2

› **Linux**

- Command line scp or sftp
- File manager, FileZilla

› **Mac OS X** (not tested)

- Command line scp or sftp
- FileZilla (Open source)




File system	Space (TB)	Quota (GB)	Backup	Shared	Cleanup	Use case
/home	26	20	yes	yes	No	Programs Code Small data sets
/data	206	250	no	yes	No	Large reused data sets
/scratch	231	-	no	yes	Yes, 30 days retention	Temporary data shared between nodes manual clean up
/local	1	-	no	per node	Yes, automatically after job	Temporary data for single node automatic clean up





› What can it do for me?





- › Account available for University staff & students
 - P, S, or F account is required
 - Undergraduate students through supervisor/teacher
 - Provide contact details and short description of planned use
- › Access through SSH protocol
 - peregrine.hpc.rug.nl (login node)
 - pg-interactive.hpc.rug.nl (interactive node)



- › Applications must be able to run under Linux
 - Compile the application yourself
 - Preinstalled applications
 - matlab, R, gromacs, ...
 - Run anywhere languages
 - Java, Python, Perl,
- › No user interaction
 - Input/output through files
- › No graphical interface





- › Long-running calculations
- › Parallel calculations
- › Many calculations



- › If you want to use commercial software a license may be necessary
- › Provided by the university
 - May be installed system wide
 - Problem with Linux versions or number of licenses
- › Provided by yourself
 - Install software in home directory
- › Other cases may be more difficult



- › Where is <insert my favourite software package>?
- › Many applications already available
- › Organized through a “module” system





- › Environment set for software when module is loaded
- › Useful commands:
 - `module avail [name]`
 - `module list`
 - `module add / load <module>`
 - `module del / remove / rm / unload <module>`
 - `module save/restore [name]`
 - `module purge`



- › Software built using toolchains:
 - goolfc:
 - › GNU compilers, OpenMPI, OpenBLAS, Lapack, FFTW, CUDA
 - ictce:
 - › Intel compilers, MKL, Intel MPI
 - Module name contains name of toolchain used to build the software
- › Dependencies automatically loaded



```
$ module avail
```

```
...
```

```
----- /software/modules/bio -----
```

```
ABYSS/1.5.2-goolfc-2.7.11-Python-2.7.9
```

```
BCFtools/1.2-goolfc-2.7.11
```

```
BEDTools/2.22.1-goolfc-2.7.11
```

```
BEDTools/2.23.0-goolfc-2.7.11 (D)
```

```
...
```

```
----- /software/modules/math -----
```

```
CPLEX/12.6.2
```

```
Eigen/3.2.4-goolfc-2.7.11
```

```
...
```

```
$ bedtools
```

```
-bash: bedtools: command not found
```

```
$ module add BEDTools/2.22.1-goolfc-2.7.11
```

```
$ bedtools --version
```

```
bedtools v2.22.1
```




```
$ module list
```

```
Currently Loaded Modules:
```

- 1) GCC/4.8.4
- 2) CUDA/6.5.14-GCC-4.8.4
- 3) gcccuda/2.7.11
- 4) OpenMPI/1.8.4-gcccuda-2.7.11
- 5) gompic/2.7.11
- 6) OpenBLAS/0.2.13-gompic-2.7.11-LAPACK-3.5.0
- 7) FFTW/3.3.4-gompic-2.7.11
- 8) ScaLAPACK/2.0.2-gompic-2.7.11-OpenBLAS-0.2.13-LAPACK-3.5.0
- 9) goolfc/2.7.11
- 10) BEDTools/2.22.1-goolfc-2.7.11

```
$ module del BEDTools
```

```
$ module list
```

```
Currently Loaded Modules:
```

- 1) GCC/4.8.4
- 2) CUDA/6.5.14-GCC-4.8.4
- 3) gcccuda/2.7.11
- 4) OpenMPI/1.8.4-gcccuda-2.7.11
- 5) gompic/2.7.11
- 6) OpenBLAS/0.2.13-gompic-2.7.11-LAPACK-3.5.0
- 7) FFTW/3.3.4-gompic-2.7.11
- 8) ScaLAPACK/2.0.2-gompic-2.7.11-OpenBLAS-0.2.13-LAPACK-3.5.0
- 9) goolfc/2.7.11



```
$ module purge
```

```
$ module list
```

```
No modules loaded
```



- › Into your own home directory:
 - Keep control over the software yourself
 - No special privileges required
 - Cannot be used by other users (unless you grant permission)

- › Into a new module:
 - Can be used by other users
 - Installation requires special privileges
 - Contact us, see “Support” slide



- › But I only see a single computer!






- › Front-end node
peregrine.hpc.rug.nl
- › Used for access to the cluster
 - Login
 - Data transfers
 - Job submission
 - Editing & Compiling programs
 - (Very) small tests





- › Interactive node:
pg-interactive.hpc.rug.nl
- › Used for access to the cluster
 - Testing and porting software
 - Data transfers
 - Job submission
 - Editing & Compiling programs
- › Shared machine, be careful about what you do!






Compute nodes

	CPU	Memory	Internal disk	Network	Accelerator
160 Standard nodes	2x Intel Xeon E5 2680v3: 24 cores @ 2.5 GHz	128 GB	1 TB	56 Gbps Infiniband + 10 Gbps ethernet	-
4 GPU nodes	2x Intel Xeon E5 2680v3: 24 cores @ 2.5 GHz	128 GB	4320 CPU cores, 23040 CUDA cores, 244 Xeon Phi cores		
2 Xeon Phi nodes	2x Intel Xeon E5 2680v3: 24 cores @ 2.5 GHz	128 GB			
7 Big memory nodes	4x Intel Xeon E7 4860v2: 48 cores @ 2.6 GHz	1024, 1536 or 2048 GB	1 TB	56 Gbps Infiniband, 10 Gbps ethernet	-
Standard desktop PC	~4 cores	~4-8GB	~1 TB	1 Gbps ethernet	Desktop GPU



- › Users write job descriptions
- › Scheduler finds matching resource
- › Scheduler tries to make optimal use of the resources
- › No resources: wait in a queue
- › Priority determined by usage of system in the recent past
- › SLURM: <http://slurm.schedmd.com>
 - › Scheduler
 - › Resource manager





	Name	Max walltime
Standard nodes	nodes (default)	10 days
Big memory	himem	10 days
GPU	gpu	3 days
Xeon Phi	phi	3 days
Short	short	30 minutes


- › Only half of the cores can be used for running long jobs (> 3 days)



- › Scheduling full nodes to jobs a waste of resources
- › Users can still request a complete node if necessary
- › Jobs run in an isolated and limited set of resources
 - Not possible to use more resources than requested
- › Jobs running on the node can still influence each other:
 - Disk space (space and bandwidth)
 - Memory (bandwidth)
 - Can sometimes be problematic



› What is a job script?





- › Job script tells the system what you want to do
 - List of steps / commands to run
- › First line should always point to the right interpreter that will run your script
 - Typically:
`#!/bin/bash`
- › Includes requirements needed to be able to run it
 - Amount of memory
 - Number of cores and/or nodes
 - Amount of time needed to complete the job



- › Can be put in job script using lines that start with `#SBATCH`
- › These lines should be at the top of the script, right after the `#!/bin/bash` line!
- › Can also be passed as command-line arguments to `sbatch` command
 - Not explained in detail here
 - One simple example:
`sbatch --time=10:00 --ntasks=4 jobscript.sh`



- › Requirements specified using various options:
 - Wall clock time
 - #SBATCH --time=*days-hh:mm:ss*
 - Number of tasks / cores:
 - #SBATCH --ntasks=*n*
 - Specific number of nodes and tasks per node:
 - #SBATCH --nodes=*m*
 - #SBATCH --ntasks-per-node=*n*
- › Specific node types using --partition option:
 - #SBATCH --partition=*himem*



- › Memory requirements can be specified using
 - #SBATCH --mem= n
 n is the total amount of memory per node (!) in megabytes
 - #SBATCH --mem-per-cpu= n
 n is the amount of memory per CPU core in megabytes
 - Suffix K or KB, M or MB, G or GB, T or TB for other units
- › Average memory limits:
 - 21.3, 32 or 42.7 GB per core on big memory nodes
 - 5.3 GB per core on all other nodes
- › Default 2000MB memory limit!
- › **Exceeding the limit will kill your application/job**



- › Also using #SBATCH lines or on the command line
- › Name of the job
 - #SBATCH --job-name=*name*
- › Name of the output file of the job
 - #SBATCH --output=*filename*
 - Default is: slurm-*<jobid>*.out
- › Email notifications
 - #SBATCH --mail-user=*user@mail.com*
 - #SBATCH --mail-type=*event1,event2,...,eventN*
 - *event* can be any of: ALL, BEGIN, END, FAIL, REQUEUE, TIME_LIMIT, TIME_LIMIT_50, TIME_LIMIT_80, TIME_LIMIT_90
 - Mail can be problematic in case of cluster problems



- › Contains Linux commands
 - cd, mkdir, etc.
- › Run some application

- › Sample script:

```
#!/bin/bash
#SBATCH --ntasks=2
#SBATCH --time=00:30:00
#SBATCH --job-name=testjob

cd $HOME/bin
./myprog
```



\$HOME: your home directory

\$USER: your username

\$SCRATCHDIR: Temporary directory created for your job on /scratch. Removed after your job has finished!

\$TMPDIR: Temporary directory created for your job on /local. Removed after your job has finished!

\$SLURM_JOB_ID: Id of job, useful for creating unique files or directories for a job

\$SLURM_JOB_NAME: Name of the job, as specified in the jobscript

...



› sbatch *jobscript*

```
$ sbatch testjob.sh  
Submitted batch job 2865
```



Job id

- › Job will start in the directory from which it was submitted
- › Your complete environment will be transferred to the job; this includes all loaded modules.
 - › Can be adjusted using:
`#SBATCH --export=VAR1,VAR2, ... (or: NONE)`



› queue

```
$ queue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
4983	nodes	testjob	p456789	PD	0:00	20	(Resources)
4984	nodes	testjob	p456789	PD	0:00	20	(Priority)
4985	nodes	testjob	p456789	PD	0:00	20	(Priority)
4986	nodes	testjob	p456789	PD	0:00	20	(Priority)
4987	nodes	testjob	p456789	PD	0:00	20	(Priority)
4978	nodes	testjob	p456789	R	0:01	20	pg-node[041-060]
4979	nodes	testjob	p456789	R	0:01	20	pg-node[061-080]
4980	nodes	testjob	p456789	R	0:01	20	pg-node[081-100]
4981	nodes	testjob	p456789	R	0:01	20	pg-node[101-120]
4982	nodes	testjob	p456789	R	0:01	20	pg-node[121-140]
4976	nodes	testjob	p456789	R	0:04	20	pg-node[001-020]
4977	nodes	testjob	p456789	R	0:04	20	pg-node[021-040]



```
$ squeue -u p456789
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
3018	nodes	hpl.128.	p456789	R	3:26	128	pg-node[001-120,122-129]

Status:

PD: pending

R: running

CA: cancelled

CG: completing

CD: completed

F: failed

NF: node failure

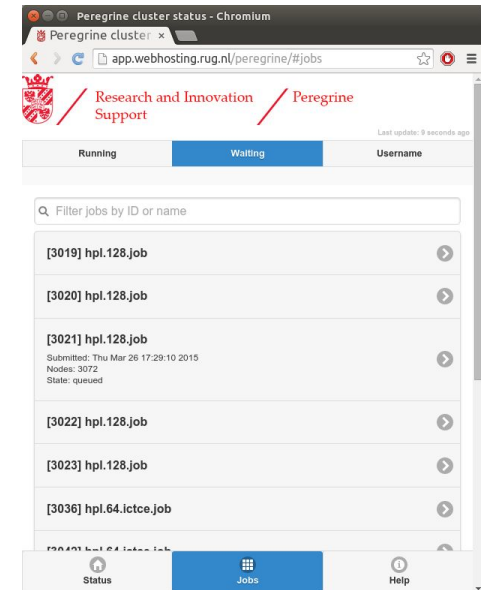
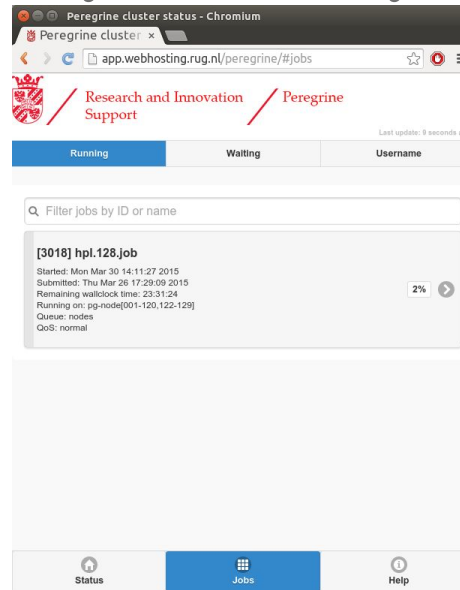

TO: timeout (reached time limit)



- › More information about a particular job:
 - › Still running or just completed:
 - › `scontrol show job <job id>`
 - › Accounting information for completed jobs:
 - › `sacct --jobs=<one or more job ids>`
 - › Add `--long` option for even more information
 - › Use `--format` option if you want specific information for your job, e.g. for memory usage:
`sacct --jobs=2781 --format=JobID,
jobname,MaxRSS`
- › Same information for running jobs:
 - › `sstat --jobs=<job id(s)>`



- › <http://app.webhosting.rug.nl>
- › Monitor cluster status and usage
- › Monitor job status, progress and information
- › Intended for smartphones, but also works on desktop
- › Also available as MyUniversity widget





- › Unless specified otherwise, output file is written to same directory as from which the job was submitted
- › *slurm-<jobid>.out*
- › Created when job starts running
- › While job is running, new output gets appended

- › If the job has state CD or has disappeared from queue, it has finished



› scancel

```
$ sbatch testjob.sh  
Submitted batch job 2870
```

```
$ squeue -u peter
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
2870	nodes	testjob	peter	R	0:03	1	pg-node021

```
$ scancel 2870
```

› Cancel multiple jobs at once:

```
$ scancel --state=PENDING --user=bob --partition=short
```





› A job script that runs R code:

```
#!/bin/bash
#SBATCH --job-name=R_job
#SBATCH --time=00:01:00
#SBATCH --ntasks=1
#SBATCH --mem=1000
#SBATCH --partition=short

pwd
module load R/3.1.2-goolfc-2.7.11-default
module list
Rscript myscript.r
```



- › A job script that runs Matlab code:

```
#!/bin/bash
#SBATCH --job-name=matlab_job
#SBATCH --time=00:02:00
#SBATCH --ntasks=1
#SBATCH --mem=1000
#SBATCH --partition=short

module load MATLAB/2014b-GCC-4.8.4
module list
matlab -nodisplay -r mycode
```


Code in file mycode.m



- › Support through CIT central service desk
 - Phone: 3232
 - E-mail: citservicedesk@rug.nl
- › Online documentation and account request form:
 - <https://redmine.hpc.rug.nl>
- › Comments and questions are always welcome




› I want to know more!





- › Introduction to Linux by Machteld Garrels:
<http://tldp.org/LDP/intro-linux/html/index.html>
- › Bash shell guide by Machteld Garrels:
<http://tldp.org/LDP/Bash-Beginners-Guide/html/index.html>
- › Linux guide: <http://www.tuxfiles.org>
- › Linux command line: <http://linuxcommand.org/>
- › Documentation and more details about SLURM:
<http://slurm.schedmd.com>
- › Manual pages for all SLURM commands:
http://slurm.schedmd.com/man_index.html





rijksuniversiteit
groningen

Peregrine team



Wietze
Albers

Fokke Dijkstra, Niels Idsinga, Ger Strikwerda, Robin Teeninga, Bob Dröge,
Laurent Jensma, Henk-Jan Zilverberg, Wim Nap



- › Hostname: peregrine.hpc.rug.nl
- › Username & password have been handed out



› sinfo

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
nodes*	up	10-00:00:0	2	down*	pg-node[006,141]
nodes*	up	10-00:00:0	2	drain	pg-node[068,111]
nodes*	up	10-00:00:0	27	resv	pg-node[005,007-015,018-021,023-035]
nodes*	up	10-00:00:0	28	alloc	pg-node[104-110,142-162]
nodes*	up	10-00:00:0	95	idle	pg-node[016-017,022,036-067,069-099,112-140]
short	up	30:00	4	mix	pg-node[100-103]
short	up	30:00	2	idle	pg-node[003-004]
gpu	up	3-00:00:00	4	comp	pg-gpu[01-04]
phi	up	3-00:00:00	2	idle	pg-phi[01-02]
himem	up	10-00:00:0	7	idle	pg-memory[01-07]



- › sview on the login node
- › See slides about running applications with a GUI

The screenshot displays the sview GUI interface. The main window, titled 'sview', has a menu bar with 'Actions', 'Options', 'Query', and 'Help'. Below the menu is a tabbed interface with 'Jobs', 'Partitions', 'Reservations', and 'Visible Tabs'. The 'Jobs' tab is active, showing a table of jobs. To the left of the table is a grid of colored squares representing the cluster's state.

JobID	Partition	UserID	Name	State	Time Running	Node Count	NodeList
3018	nodes	fokke	hpl.128.job	RUNNING	00:39:55	128	pg-node[001-120,
3021	nodes	fokke	hpl.128.job	PENDING	00:00:00	128	waiting...
3020	nodes	fokke	hpl.128.job	PENDING	00:00:00	128	waiting...
3019	nodes	fokke	hpl.128.job	PENDING	00:00:00	128	
3023	nodes	fokke	hpl.128.job	PENDING	00:00:00	128	
3022	nodes	fokke	hpl.128.job	PENDING	00:00:00	128	
3043	nodes	fokke	hpl.64.ictce.job	PENDING	00:00:00	64	
3042	nodes	fokke	hpl.64.ictce.job	PENDING	00:00:00	64	
3044	nodes	fokke	hpl.64.ictce.job	PENDING	00:00:00	64	
3046	nodes	fokke	hpl.64.ictce.job	PENDING	00:00:00	64	
3045	nodes	fokke	hpl.64.ictce.job	PENDING	00:00:00	64	
3049	nodes	fokke	hpl.64.ictce.job	PENDING	00:00:00	64	

A pop-up window titled 'Full info for job 3018' is open, showing detailed information for the selected job. It includes a smaller version of the cluster grid and a list of job specifications.

Property	Value
Corespec	65534
CPU Count	3072
CPUs Max	0
CPUs Min	3072
CPUs per Task	1
Dependency	
Derived Exit Code	0:0
Exit Code	0:0
Features	
Gres	
GroupID	beheer
JobID	3018
Licenses	
Min CPUs Per Node	24
Min Memory	120G Per Node

The pop-up window has buttons for 'Refresh', 'Close', and 'Close All Popups' at the bottom.